

1 Sources

BAS TAXI 2.5 (real speech)

German dispatcher ↔ English client
86 sessions, 640 usable turns
human translations, 8 kHz phone
push-to-talk: no timing, no overlap

CST-Bench-Syn (evaluation only)

XDailyDialog, BConTrasT dialogues
EN↔DE: human translations
KO↔EN: Korean side machine-
translated, consensus-checked
TTS speech

planned

2 Common schema + audio

manifest.jsonl

per turn: speaker, language,
transcript, reference translation
8 → 16 kHz (TAXI)
energy-based silence trim,
0.1 s margin kept

3 Deterministic session renderer

turns joined in real dialogue order;
seed = SHA-1(session id) + config seed

L0 natural

gap ~ $N(0.2, 0.25^2)$ s, clipped to [0.05, 1.0]

L0 mediated

gap ~ $U[1, 3]$ s

L1 (planned)

next turn starts 0.2–0.8 s before the end

outputs per session

DE EN

mix.wav

single-channel input

2ch.wav

per speaker (analysis)

timeline.json

turn times, references

4 System under test + scorer

system hears mix.wav only and emits
committed pieces (time, language, text)

cstbench eval

per direction: min-edit-distance
resegmentation onto reference turns

quality BLEU, chrF, COMET*

lag StreamLAAL, EndOffset,
switch latency*

errors empty turns,
wrong direction*

aux. WER*, cpWER*, DER*

* not yet in cstbench v0.1

Released: code (Apache-2.0), configs and checksums only; TAXI audio and text are never redistributed (rebuild verified: 86/86 timeline digests).