



(a)




(b)



(c)

prefix (input only): ... language pair English–German <DELAY_2> <TDELAY_3> then, per 80 ms chunk k: audio embedding, speaker-tagged transcript, target-tagged translation, <next>



a_k audio embedding of chunk k (input, no loss) <next> <NEXT_AUDIO>: end of chunk, read 80 ms more word transcript token (tint = speaker's language)

[A] speaker tag (arrival order) [<-de>] target-language tag Taxi translation token — underline = training target

dashed = planned, not yet built